

```
In [7]: # Traffic Sign Detection Using CNN
```

```
In [8]: #initializations  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import tensorflow as tf  
import cv2  
from PIL import Image  
import os  
from tensorflow import keras  
from keras.datasets import mnist  
from tensorflow.keras.layers import Conv2D, MaxPooling2D  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Dropout, Flatten  
from tensorflow.keras.optimizers import SGD  
from tensorflow.keras import backend as K
```

Using TensorFlow backend.

```

In [9]: ## Reading images and putting in numpy
data = []
labels = []

height = 30
width = 30
channels = 3
classes = 43
n_inputs = height * width * channels

for i in range(classes) :
    path = "C:/Users/HP/Desktop/TrafficSign/Train/{0}/".format(i)
    print(path)
    Class=os.listdir(path)
    for a in Class:
        try:
            image=cv2.imread(path+a)
            image_from_array = Image.fromarray(image, 'RGB')
            size_image = image_from_array.resize((height, width))
            data.append(np.array(size_image))
            labels.append(i)
        except AttributeError:
            print(" ")

Cells=np.array(data)
labels=np.array(labels)
##Randomize the order of the input images because accuracy was too high and valid
s=np.arange(Cells.shape[0])
np.random.seed(43)
np.random.shuffle(s)

Cells=Cells[s]
labels=labels[s]

```

```

C:/Users/HP/Desktop/TrafficSign/Train/0/
C:/Users/HP/Desktop/TrafficSign/Train/1/
C:/Users/HP/Desktop/TrafficSign/Train/2/
C:/Users/HP/Desktop/TrafficSign/Train/3/
C:/Users/HP/Desktop/TrafficSign/Train/4/
C:/Users/HP/Desktop/TrafficSign/Train/5/
C:/Users/HP/Desktop/TrafficSign/Train/6/
C:/Users/HP/Desktop/TrafficSign/Train/7/
C:/Users/HP/Desktop/TrafficSign/Train/8/
C:/Users/HP/Desktop/TrafficSign/Train/9/
C:/Users/HP/Desktop/TrafficSign/Train/10/
C:/Users/HP/Desktop/TrafficSign/Train/11/
C:/Users/HP/Desktop/TrafficSign/Train/12/
C:/Users/HP/Desktop/TrafficSign/Train/13/
C:/Users/HP/Desktop/TrafficSign/Train/14/
C:/Users/HP/Desktop/TrafficSign/Train/15/
C:/Users/HP/Desktop/TrafficSign/Train/16/
C:/Users/HP/Desktop/TrafficSign/Train/17/
C:/Users/HP/Desktop/TrafficSign/Train/18/
C:/Users/HP/Desktop/TrafficSign/Train/19/
C:/Users/HP/Desktop/TrafficSign/Train/20/
C:/Users/HP/Desktop/TrafficSign/Train/21/

```

```
C:/Users/HP/Desktop/TrafficSign/Train/22/  
C:/Users/HP/Desktop/TrafficSign/Train/23/  
C:/Users/HP/Desktop/TrafficSign/Train/24/  
C:/Users/HP/Desktop/TrafficSign/Train/25/  
C:/Users/HP/Desktop/TrafficSign/Train/26/  
C:/Users/HP/Desktop/TrafficSign/Train/27/  
C:/Users/HP/Desktop/TrafficSign/Train/28/  
C:/Users/HP/Desktop/TrafficSign/Train/29/  
C:/Users/HP/Desktop/TrafficSign/Train/30/  
C:/Users/HP/Desktop/TrafficSign/Train/31/  
C:/Users/HP/Desktop/TrafficSign/Train/32/  
C:/Users/HP/Desktop/TrafficSign/Train/33/  
C:/Users/HP/Desktop/TrafficSign/Train/34/  
C:/Users/HP/Desktop/TrafficSign/Train/35/  
C:/Users/HP/Desktop/TrafficSign/Train/36/  
C:/Users/HP/Desktop/TrafficSign/Train/37/  
C:/Users/HP/Desktop/TrafficSign/Train/38/  
C:/Users/HP/Desktop/TrafficSign/Train/39/  
C:/Users/HP/Desktop/TrafficSign/Train/40/  
C:/Users/HP/Desktop/TrafficSign/Train/41/  
C:/Users/HP/Desktop/TrafficSign/Train/42/
```

In [10]: *## Train Test Split 80-20*

```
(x_train,x_test) = Cells[(int)(0.2*len(labels)):],Cells[::(int)(0.2*len(labels))]  
(y_train,y_test) = labels[(int)(0.2*len(labels)):],labels[::(int)(0.2*len(labels))]
```

In [11]: *## Train-test Shapes*

```
print(x_train.shape) #train images  
print(y_train.shape) #train labels  
print(x_test.shape) #test images  
print(y_test.shape) #test labels
```

```
(31368, 30, 30, 3)  
(31368,)   
(7841, 30, 30, 3)  
(7841,)
```

In [12]: *## Setting input image shape*

```
#train[0] gives 1st dimension and train[1] gives second dimension  
img_rows = x_train[0].shape[0]  
img_cols = x_train[1].shape[0]  
  
#storing shape of a single image  
input_shape = (img_rows, img_cols, 3)  
print(input_shape)
```

```
(30, 30, 3)
```

```
In [13]: ## Preprocessing of image data

#keras want float32 data type
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
#normalization of image data ... as images range from 0 to 255 and normalize them
x_train = x_train / 255
x_test = x_test / 255
```

```
In [14]: #hot one encoding

from keras.utils import np_utils

y_train = np_utils.to_categorical(y_train, classes)
#print(y_train)
y_test = np_utils.to_categorical(y_test, classes)
#print(y_test)
```

```

In [15]: #model creation
model = Sequential()

#conv layer1
model.add(Conv2D(32, kernel_size=(3,3), activation='relu',input_shape = input_shape))
model.add(MaxPooling2D(pool_size = (2,2)))

#conv layer2
model.add(Conv2D(64, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

#conv layer3
model.add(Conv2D(128, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size = (2,2)))

#flatten
model.add(Flatten())

#dense layer 1
model.add(Dense(256,activation = 'relu'))
model.add(Dropout(0.25))

#dense layer 2
model.add(Dense(classes, activation = 'softmax'))
#compiling model
model.compile(loss = 'categorical_crossentropy',
              optimizer = 'adam',
              metrics = ['accuracy'])
print(model.summary())

#training begins here
#verbose is how much information we wanna see in training
batch_size = 32
epochs = 20
#batch size 16 for large images and 32 for small
history = model.fit(x_train, y_train,
                    batch_size = batch_size,
                    epochs = epochs,
                    verbose = 1,
                    validation_data = (x_test,y_test))

#evaluate model
score = model.evaluate(x_test, y_test, verbose = 0)
print('Test loss',score[0])
print('Test accuracy',score[1])

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	896

max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout (Dropout)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 128)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 43)	11051
=====		
Total params: 235,627		
Trainable params: 235,627		
Non-trainable params: 0		

None

Train on 31368 samples, validate on 7841 samples

Epoch 1/20

31368/31368 [=====] - 31s 974us/sample - loss: 1.3730
- accuracy: 0.6057 - val_loss: 0.2610 - val_accuracy: 0.9168

Epoch 2/20

31368/31368 [=====] - 30s 945us/sample - loss: 0.2257
- accuracy: 0.9322 - val_loss: 0.1099 - val_accuracy: 0.9749

Epoch 3/20

31368/31368 [=====] - 30s 954us/sample - loss: 0.1179
- accuracy: 0.9649 - val_loss: 0.0748 - val_accuracy: 0.9790

Epoch 4/20

31368/31368 [=====] - 30s 959us/sample - loss: 0.0820
- accuracy: 0.9741 - val_loss: 0.0454 - val_accuracy: 0.9884

Epoch 5/20

31368/31368 [=====] - 32s 1ms/sample - loss: 0.0590 -
accuracy: 0.9820 - val_loss: 0.0482 - val_accuracy: 0.9879

Epoch 6/20

31368/31368 [=====] - 35s 1ms/sample - loss: 0.0567 -
accuracy: 0.9824 - val_loss: 0.0275 - val_accuracy: 0.9939

Epoch 7/20

31368/31368 [=====] - 36s 1ms/sample - loss: 0.0443 -
accuracy: 0.9861 - val_loss: 0.0341 - val_accuracy: 0.9913

Epoch 8/20

31368/31368 [=====] - 37s 1ms/sample - loss: 0.0388 -
accuracy: 0.9880 - val_loss: 0.0273 - val_accuracy: 0.9939

Epoch 9/20

31368/31368 [=====] - 36s 1ms/sample - loss: 0.0355 -
accuracy: 0.9886 - val_loss: 0.0347 - val_accuracy: 0.9925

Epoch 10/20

31368/31368 [=====] - 36s 1ms/sample - loss: 0.0343 -
accuracy: 0.9904 - val_loss: 0.0382 - val_accuracy: 0.9895

Epoch 11/20

```
31368/31368 [=====] - 35s 1ms/sample - loss: 0.0273 -  
accuracy: 0.9914 - val_loss: 0.0268 - val_accuracy: 0.9936  
Epoch 12/20  
31368/31368 [=====] - 34s 1ms/sample - loss: 0.0267 -  
accuracy: 0.9911 - val_loss: 0.0203 - val_accuracy: 0.9955  
Epoch 13/20  
31368/31368 [=====] - 34s 1ms/sample - loss: 0.0295 -  
accuracy: 0.9911 - val_loss: 0.0223 - val_accuracy: 0.9957  
Epoch 14/20  
31368/31368 [=====] - 34s 1ms/sample - loss: 0.0261 -  
accuracy: 0.9917 - val_loss: 0.0308 - val_accuracy: 0.9935  
Epoch 15/20  
31368/31368 [=====] - 35s 1ms/sample - loss: 0.0233 -  
accuracy: 0.9930 - val_loss: 0.0230 - val_accuracy: 0.9950  
Epoch 16/20  
31368/31368 [=====] - 36s 1ms/sample - loss: 0.0250 -  
accuracy: 0.9924 - val_loss: 0.0278 - val_accuracy: 0.9926  
Epoch 17/20  
31368/31368 [=====] - 37s 1ms/sample - loss: 0.0211 -  
accuracy: 0.9936 - val_loss: 0.0302 - val_accuracy: 0.9925  
Epoch 18/20  
31368/31368 [=====] - 37s 1ms/sample - loss: 0.0227 -  
accuracy: 0.9932 - val_loss: 0.0203 - val_accuracy: 0.9941  
Epoch 19/20  
31368/31368 [=====] - 38s 1ms/sample - loss: 0.0215 -  
accuracy: 0.9932 - val_loss: 0.0283 - val_accuracy: 0.9938  
Epoch 20/20  
31368/31368 [=====] - 35s 1ms/sample - loss: 0.0224 -  
accuracy: 0.9936 - val_loss: 0.0179 - val_accuracy: 0.9957  
Test loss 0.017889902631156914  
Test accuracy 0.9956638
```

```

In [16]: #plot the loss charts
history_dict = history.history


loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(loss_values)+1)

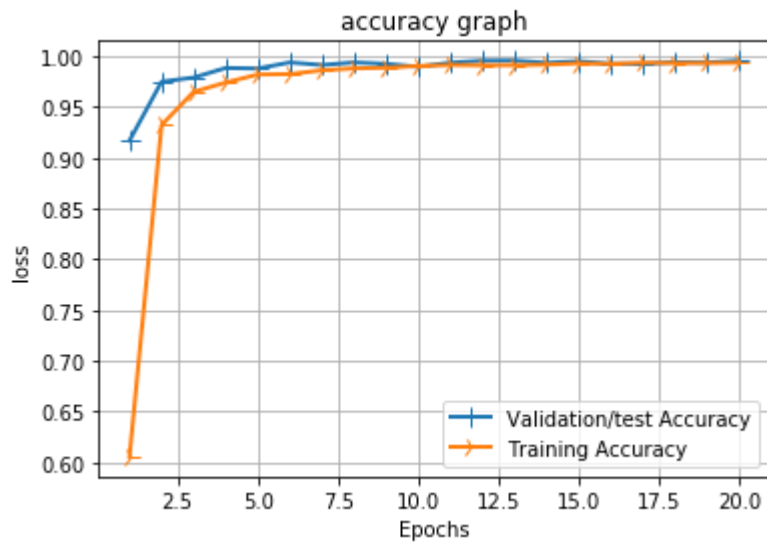
line1 = plt.plot(epochs,val_loss_values, label = 'Validation/test loss')
line2 = plt.plot(epochs,loss_values, label = 'Training loss')
plt.setp(line1, linewidth = 2.0, marker = '+', markersize = 10.0 )
plt.setp(line2, linewidth = 2.0, marker = '4', markersize = 10.0 )
plt.title('loss graph')
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.grid(True)
plt.legend()
plt.show()

acc_values = history_dict['accuracy']
val_acc_values = history_dict['val_accuracy']

line1 = plt.plot(epochs,val_acc_values, label = 'Validation/test Accuracy')
line2 = plt.plot(epochs,acc_values, label = 'Training Accuracy')
plt.setp(line1, linewidth = 2.0, marker = '+', markersize = 10.0 )
plt.setp(line2, linewidth = 2.0, marker = '4', markersize = 10.0 )
plt.title('accuracy graph')
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.grid(True)
plt.legend()
plt.show()

```





```
In [17]: #save the model
from keras.models import model_from_json
model.save("C:/Users/HP/Desktop/TrafficSign/cnn_signal.h5")
print('model saved')
```

model saved

```
In [18]: #Load model
from tensorflow.keras.models import load_model
classifier = load_model("C:/Users/HP/Desktop/TrafficSign/cnn_signal.h5")
```

```
In [19]: #Saving the history File of our model
import pickle

pickle_out = open("trafficSign_pickle","wb")
pickle.dump(history.history, pickle_out)
pickle_out.close()
```

In [20]: *#To load the save history of model*

```
pickle_in = open("trafficSign_pickle","rb")
saved_history = pickle.load(pickle_in)
print(saved_history)
```

```
{'loss': [1.3729721551204754, 0.22566022191674573, 0.11794098528390391, 0.08201463327231728, 0.0589599325390496, 0.056652773496607814, 0.044309441444429096, 0.03876699338290267, 0.03550584499654482, 0.03428600539905851, 0.027335879033789118, 0.026672158697717382, 0.029452142662001104, 0.026137571665771287, 0.023337168610882073, 0.024977875898564407, 0.021051364324426232, 0.02266189936173432, 0.021489173710362245, 0.022433864524114167], 'accuracy': [0.60568094, 0.9322239, 0.96490055, 0.9740819, 0.9819561, 0.98237056, 0.9861005, 0.98801327, 0.98861897, 0.99040425, 0.9913606, 0.9910737, 0.99113744, 0.9917432, 0.9929546, 0.9924445, 0.9935603, 0.9931778, 0.99320966, 0.9935922], 'val_loss': [0.2610037137431677, 0.10991166436332545, 0.07476574632861425, 0.04535789232614401, 0.04816794988281464, 0.027491415433260915, 0.03408292307128107, 0.02727517366098761, 0.03471781250152566, 0.03822880742013202, 0.026828537457087392, 0.020307183260564285, 0.02225077899036963, 0.030787936957357487, 0.023046976764631088, 0.027837193084593752, 0.030226810949427106, 0.020292335452397316, 0.02829721574097263, 0.017889902631156914], 'val_accuracy': [0.91684735, 0.9748756, 0.97895676, 0.9883943, 0.9878842, 0.9938783, 0.99132764, 0.9938783, 0.99247545, 0.9895421, 0.99362326, 0.99553627, 0.9956638, 0.9934957, 0.9950262, 0.992603, 0.99247545, 0.9941334, 0.9937508, 0.9956638]}
```

```
In [26]: import sklearn
import scipy
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

y_pred = classifier.predict_classes(x_test)

#classification report
print(classification_report(np.argmax(y_test,axis=1), y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	39
1	1.00	0.99	0.99	451
2	0.99	0.99	0.99	481
3	0.97	1.00	0.98	274
4	1.00	1.00	1.00	411
5	0.99	0.97	0.98	346
6	0.98	1.00	0.99	91
7	1.00	0.99	1.00	318
8	1.00	1.00	1.00	280
9	1.00	1.00	1.00	291
10	1.00	1.00	1.00	379
11	1.00	1.00	1.00	249
12	1.00	1.00	1.00	366
13	1.00	1.00	1.00	449
14	1.00	0.99	1.00	145
15	0.99	1.00	1.00	146
16	1.00	1.00	1.00	95
17	1.00	1.00	1.00	208
18	0.98	1.00	0.99	249
19	0.98	1.00	0.99	42
20	1.00	0.99	0.99	73
21	1.00	0.98	0.99	63
22	1.00	1.00	1.00	90
23	1.00	1.00	1.00	93
24	1.00	0.98	0.99	64
25	1.00	1.00	1.00	304
26	1.00	0.97	0.99	117
27	1.00	0.96	0.98	57
28	1.00	1.00	1.00	95
29	1.00	1.00	1.00	57
30	1.00	0.99	0.99	98
31	0.99	1.00	1.00	168
32	1.00	1.00	1.00	55
33	1.00	1.00	1.00	143
34	1.00	1.00	1.00	78
35	1.00	1.00	1.00	241
36	1.00	1.00	1.00	68
37	1.00	1.00	1.00	37
38	1.00	1.00	1.00	411
39	1.00	1.00	1.00	53
40	0.99	1.00	0.99	71
41	1.00	1.00	1.00	45
42	1.00	1.00	1.00	50

accuracy			1.00	7841
macro avg	1.00	1.00	1.00	7841
weighted avg	1.00	1.00	1.00	7841



```
In [27]: #Confusion Matrix
print(confusion_matrix(np.argmax(y_test,axis=1),y_pred))
```

```
[[ 39   0   0 ...   0   0   0]
 [   0 448   1 ...   0   0   0]
 [   0   2 476 ...   0   0   0]
 ...
 [   0   0   0 ...  71   0   0]
 [   0   0   0 ...   0  45   0]
 [   0   0   0 ...   0   0  50]]
```

```
In [ ]:
```